

Kernel Modules

Kernel Module

- Portion of kernel that can be dynamically loaded and unloaded
- Examples
 - USB drivers
 - File system drivers
 - Disk drivers
 - Cryptographic libraries

Why modules?

- Not every machine needs the same modules
 - Different machines use different drivers
- Load only the components you need
 - Smaller system footprint
 - Quicker boot time
- Dynamically load modules for new devices
 - New USB, camera, printer
 - Changing graphics card, motherboard, file system

Kernel Logistics

- Source code is stored in /usr/src/
- Kernel image gets installed to /boot/vmlinux-<kernel-name>

Kernel Programming

- Kernel modules are event-driven
 - Register functions
 - Wait for requests from user-space and service them
 - Server/client model
- No standard C library → kernel libraries instead
- No floating point support
- Crashes/deadlocks in a module can cause the entire kernel to crash
 - Requires system-wide reboot

Kernel Headers

- `#include <linux/init.h>`
 - Module stuff
- `#include <linux/module.h>`
 - Module stuff
- `#include <asm/semaphore.h>`
 - Locks
- `#include <linux/list.h>`
 - Linked lists
- `#include <linux/string.h>`
 - String functions
- Can find others at <https://elixir.bootlin.com/>
 - `/include/linux/`

printk

- Similar to printf but prints to the kernel log
- Takes log level and format string as parameters

```
printk(KERN_INFO "hello %s\n", str_var);
```
- Note there is no comma between the log level and string argument
- Kernel log viewable through
 - > cat /var/log/kern.log
 - > dmesg
 - > tail -f /var/log/kern.log (to watch in real-time)

printk

- Log levels:
 - KERN_EMERG Emergency condition, kernel likely crashed
 - KERN_ALERT Alert that requires immediate attention
 - KERN_CRIT Critical error message
 - KERN_ERR Error message
 - KERN_WARNING Warning message
 - KERN_NOTICE Normal, but noteworthy message
 - KERN_INFO Informational message
 - KERN_DEBUG Debug message

kmalloc

- `void * kmalloc (size_t size, gfp_t flags);`
- `char_ptr = kmalloc (sizeof(char) * 20, __GFP_RECLAIM)`
- Remember to restrict kernel memory allocation
 - Can block important functions
 - Can crash kernel if improperly handled
 - Kernel has limited access to memory

kmalloc flags

- `__GFP_RECLAIM` Allocator can sleep
- `__GFP_HIGH` Allocator can access emergency pools
- `__GFP_IO` Allocator can start disk I/O
- `__GFP_FS` Allocator can start filesystem I/O
- `__GFP_COLD` Allocator should use cache cold pages
- `__GFP_NOWARN` Allocator will not print failure warnings
- `__GFP_REPEAT` Allocator will repeat if it fails (can still fail)
- `__GFP_NOFAIL` Allocator will repeat if it fails (can not fail)
- `__GFP_NORETRY` Allocator will never retry if it fails
- `__GFP_NO_GROW` Used by the slab
- `__GFP_COMP` Used by hugetlb

Hello World Kernel Module

- Install development tools:
 - > `sudo apt-get install build-essential linux-headers-`uname -r``
- Create a folder in `/usr/src/` called `hello`
- Create the kernel module program (`/usr/src/hello/hello.c`)
- Compile
- Insert/load the kernel module
- Remove/unload the kernel module
- Tutorial: <https://blog.sourcerer.io/writing-a-simple-linux-kernel-module-d9dc3762c234>

hello.c

```
#include <linux/init.h>
#include <linux/module.h>

MODULE_LICENSE("Dual BSD/GPL");

static int hello_init(void)
{
    printk(KERN_ALERT "Hello, world!\n");
    return 0;
}

static void hello_exit(void)
{
    printk(KERN_ALERT "Goodbye, world!\n");
    return;
}

module_init(hello_init);
module_exit(hello_exit);
```

hello.c

```
#include <linux/init.h>
#include <linux/module.h>
```

Linux module headers

```
MODULE_LICENSE("Dual BSD/GPL");
```

```
static int hello_init(void)
{
    printk(KERN_ALERT "Hello, world!\n");
    return 0;
}
```

```
static void hello_exit(void)
{
    printk(KERN_ALERT "Goodbye, world!\n");
    return;
}
```

```
module_init(hello_init);
module_exit(hello_exit);
```

hello.c

```
#include <linux/init.h>
#include <linux/module.h>
```

```
MODULE_LICENSE("Dual BSD/GPL");
```



License declaration

```
static int hello_init(void)
{
    printk(KERN_ALERT "Hello, world!\n");
    return 0;
}
```

```
static void hello_exit(void)
{
    printk(KERN_ALERT "Goodbye, world!\n");
    return;
}
```

```
module_init(hello_init);
module_exit(hello_exit);
```

hello.c

```
#include <linux/init.h>
```

```
#include <linux/module.h>
```

```
MODULE_LICENSE("Dual BSD/GPL");
```

```
static int hello_init(void)
```

```
{  
    printk(KERN_ALERT "Hello, world!\n");  
    return 0;  
}
```

Initialization function

```
static void hello_exit(void)
```

```
{  
    printk(KERN_ALERT "Goodbye, world!\n");  
    return;  
}
```

```
module_init(hello_init);
```

```
module_exit(hello_exit);
```

Run when module is loaded

hello.c

```
#include <linux/init.h>
#include <linux/module.h>
```

```
MODULE_LICENSE("Dual BSD/GPL");
```

```
static int hello_init(void)
{
    printk(KERN_ALERT "Hello, world!\n");
    return 0;
}
```

```
static void hello_exit(void)
{
    printk(KERN_ALERT "Goodbye, world!\n");
    return;
}
```



Exit function

```
module_init(hello_init);
module_exit(hello_exit);
```



Run when module is unloaded

Kernel Makefile

```
obj-m += hello.o
```

```
all:
```

```
    make -C /lib/modules/<kernel-version>/build M=/usr/src/hello modules
```

```
clean:
```

```
    make -C /lib/modules/<kernel-version>/build M=/usr/src/hello clean
```

(<kernel-version> can be found using `uname -r`)

(place Makefile in /usr/src/hello)

Compilation

- To compile:
`/usr/src/hello > sudo make`
- This creates kernel object `hello.ko`

Loading the Module

- Load the module:

```
/usr/src/hello > sudo insmod hello.ko
```

- Check the module is loaded:

```
> lsmod | grep hello
```

- Check the kernel log:

```
> dmesg
```

- Unload the kernel module:

```
> sudo rmmod hello
```