

System Calls

What are they?

- Standard interface to allow the kernel to safely handle user requests. Example requests:
 - Read from hardware
 - Spawn a new process
 - Get current time
 - Create shared memory
- Message passing technique between:
 - OS kernel (server)
 - User (client)

Executing System Calls

- User program issues call
- Core kernel looks up the call in the syscall table
- *Kernel module* handles syscall action
- Module returns result of system call
- Core kernel forwards results to the user

What if the module is not loaded?

- User program issues call
 - Core kernel looks up the call in the syscall table
 - Kernel module isn't loaded to handle action
 - ...
-
- Where does call go?

System Call Wrappers

- Have a system call wrapper to handle this scenario.
- Wrapper calls the system call handler function if the module is loaded
- Else, it returns an error
- Uses a function pointer to point the system call handler function
- Add a system call wrapper for each system call you add.

Adding System Calls

- Files to add:
 - `/usr/src/test_kernel/SystemCalls/test_call.c`
 - `/usr/src/test_kernel/SystemCalls/Makefile`
 - `/usr/src/test_kernel/SystemCalls/syscallModule.c`
- Files to modify:
 - `/usr/src/test_kernel/arch/x86/entry/syscalls/syscall_64.tbl`
 - `/usr/src/test_kernel/include/linux/syscalls.h`
 - `/usr/src/test_kernel/Makefile`

SystemCalls/test_call.c

```
#include <linux/linkage.h>
#include <linux/kernel.h>
#include <linux/module.h>
#include <linux/syscalls.h>

/* System call stub */
long (*STUB_test_call)(int) = NULL;
EXPORT_SYMBOL(STUB_test_call);

/* System call wrapper */
SYSCALL_DEFINE1(test_call, int, test_int) {
    printk(KERN_NOTICE "Inside SYSCALL_DEFINE1 block.
        %s: Your int is %d\n", __FUNCTION__, test_int);
    if (STUB_test_call != NULL)
        return STUB_test_call(test_int);
    else
        return -ENOSYS;
}
```

- Creates syscall pointer.
- Exports the pointer so that the system call module can access it.
- Define syscall wrapper.

SystemCalls/test_call.c

```
#include <linux/linkage.h>
#include <linux/kernel.h>
#include <linux/module.h>
#include <linux/syscalls.h>

/* System call stub */
long (*STUB_test_call)(int) = NULL;
EXPORT_SYMBOL(STUB_test_call);

/* System call wrapper */
SYSCALL_DEFINE1(test_call, int, test_int) {
    printk(KERN_NOTICE "Inside SYSCALL_DEFINE1 block.
        %s: Your int is %d\n", __FUNCTION__, test_int);
    if (STUB_test_call != NULL)
        return STUB_test_call(test_int);
    else
        return -ENOSYS;
}
```

syscall pointer (function pointer)

Exports syscall pointer so that the handler module can access it

Wrapper function

Execute if defined

Return error if not defined.

SYSCALL_DEFINE

- SYSCALL_DEFINE is a macro that generates the proper system call definition with the appropriate type of arguments.
- In SYSCALL_DEFINE1, 1 means the system call will take 1 argument.
- *SYSCALL_DEFINE1(test_call, int, test_int)* creates a system call named *sys_test_call*, which takes one argument *test_int* which is of type *int*. (Notice the comma between *int* and *test_int*).
- SYSCALL_DEFINE0 to SYSCALL_DEFINE6 are defined, so you can pass from zero to maximum six arguments to a system call.

```

#include <linux/init.h>
#include <linux/kernel.h>
#include <linux/module.h>
#include <linux/linkage.h>
MODULE_LICENSE("GPL");

extern long (*STUB_test_call)(int);
long my_test_call(int test) {
    printk(KERN_NOTICE "%s: Your int is %d\n",
        __FUNCTION__, test);
    return test;
}

static int hello_init(void) {
    STUB_test_call = my_test_call;
    return 0;
}
module_init(hello_init);

static void hello_exit(void) {
    STUB_test_call = NULL;
}
module_exit(hello_exit);

```

SystemCalls/syscallModule.c

- Holds module code.
- Registers syscall pointer to the proper syscall handler.
- Implements syscall behavior.

SystemCalls/syscallModule.c

```
#include <linux/init.h>
#include <linux/kernel.h>
#include <linux/module.h>
#include <linux/linkage.h>
MODULE_LICENSE("GPL");

extern long (*STUB_test_call) (int);

long my_test_call(int test) {
    printk(KERN_NOTICE "%s: Your int is %d\n",
        __FUNCTION__, test);
    return test;
}

static int hello_init(void) {
    STUB_test_call = my_test_call;
    return 0;
}

module_init(hello_init);

static void hello_exit(void) {
    STUB_test_call = NULL;
}

module_exit(hello_exit);
```

Get access to the syscall pointer

Actual system call handler.

Assigns the syscall pointer to the syscall handler function when the module is loaded.

Empties the syscall pointer when the module is unloaded.

SystemCalls/Makefile

obj-y := test_call.o

obj-m := syscallModule.o

PWD := \$(shell pwd)

KDIR := /lib/modules/`uname -r`/build

default:

\$(MAKE) -C \$(KDIR) SUBDIRS= \$(PWD) modules

clean:

rm -f *.o *.ko *.mod.* Module.* modules.*

arch/x86/entry/syscalls/syscall_64.tbl

```
GNU nano 2.9.3 arch/x86/entry/syscalls/syscall_64.tbl Modified
316 common renameat2 __x64_sys_renameat2
317 common seccomp __x64_sys_seccomp
318 common getrandom __x64_sys_getrandom
319 common memfd_create __x64_sys_memfd_create
320 common kexec_file_load __x64_sys_kexec_file_load
321 common bpf __x64_sys_bpf
322 64 execveat __x64_sys_execveat/ptregs
323 common userfaultfd __x64_sys_userfaultfd
324 common membarrier __x64_sys_membarrier
325 common mlock2 __x64_sys_mlock2
326 common copy_file_range __x64_sys_copy_file_range
327 64 preadv2 __x64_sys_preadv2
328 64 pwritev2 __x64_sys_pwritev2
329 common pkey_mprotect __x64_sys_pkey_mprotect
330 common pkey_alloc __x64_sys_pkey_alloc
331 common pkey_free __x64_sys_pkey_free
332 common statx __x64_sys_statx
333 common io_pgetevents __x64_sys_io_pgetevents
334 common rseq __x64_sys_rseq
335 common test_call __x64_sys_test_call
#
# x32-specific system call numbers start at 512 to avoid cache impact
# for native 64-bit operation. The __x32_compat_sys stubs are created
# on-the-fly for compat_sys_*( ) compatibility system calls if X86_X32
# is defined.
#
512 x32 rt_sigaction __x32_compat_sys_rt_sigaction
513 x32 rt_sigreturn sys32_x32_rt_sigreturn
514 x32 ioctl __x32_compat_sys_ioctl
515 x32 readv __x32_compat_sys_readv
516 x32 writev __x32_compat_sys_writev
```

- Here, we add our system call into the table.
- Up to 334 was already there, so we take 335.
- Add the line:
335<tab>common<tab>test_call<tab>__x64_sys_test_call
- Notice how the name becomes __x64_sys_test_call from our simple test_call

Include/linux/syscalls.h

```
GNU nano 2.9.3      include/linux/syscalls.h

static inline long ksys_truncate(const char __user *pathname, loff_t length)
{
    return do_sys_truncate(pathname, length);
}

static inline unsigned int ksys_personality(unsigned int personality)
{
    unsigned int old = current->personality;

    if (personality != 0xffffffff)
        set_personality(personality);

    return old;
}

asmlinkage long sys_test_call(int);

#endif
```

- End of document
- Define the system call prototype
- Notice this time we write `sys_test_call`, instead of `test_call` or `_____x64_sys_test_call`

Makefile

```
GNU nano 2.9.3 Makefile Modified

HOST_LIBELF_LIBS = $(shell pkg-config libelf --libs 2>/dev/null || echo -lelf)

ifdef CONFIG_STACK_VALIDATION
    has_libelf := $(call try-run,\
        echo "int main() {}" | $(HOSTCC) -xc -o /dev/null $(HOST_LIBELF_LIBS) -,1,0)
    ifeq ($(has_libelf),1)
        objtool_target := tools/objtool FORCE
    else
        SKIP_STACK_VALIDATION := 1
        export SKIP_STACK_VALIDATION
    endif
endif

ifeq ($(KBUILD_EXTMOD),)
core-y += kernel/ certs/ mm/ fs/ ipc/ security/ crypto/ block/ SystemCalls/

vmlinux-dirs := $(patsubst %/,%, $(filter %/, $(init-y) $(init-m) \
    $(core-y) $(core-m) $(drivers-y) $(drivers-m) \
    $(net-y) $(net-m) $(libs-y) $(libs-m) $(virt-y)))

vmlinux-alldirs := $(sort $(vmlinux-dirs) $(patsubst %/,%, $(filter %/, \
    $(init-) $(core-) $(drivers-) $(net-) $(libs-) $(virt-))))

init-y := $(patsubst %/, %/built-in.a, $(init-y))
core-y := $(patsubst %/, %/built-in.a, $(core-y))
drivers-y := $(patsubst %/, %/built-in.a, $(drivers-y))
```

- Search for the second occurrence of *corey*
- Add the *SystemCalls* directory to the list. These are the directories that have files to be built directly into the kernel.

User space program

```
1 #define _GNU_SOURCE
2 #include <stdio.h>
3 #include <stdlib.h>
4 #include <unistd.h>
5 #include <sys/syscall.h>
6 #define __NR_TEST_CALL 335
7
8 int test_call(int test) {
9     return syscall(__NR_TEST_CALL, test);
10 }
11
12 int main(int argc, char **argv) {
13     if (argc != 2) {
14         printf("wrong number of args\n");
15         return -1;
16     }
17
18     int test = atoi(argv[1]);
19     long ret = test_call(test);
20
21     if (ret < 0)
22         perror("system call error");
23     else
24         printf("Function successful. passed in: %d,\n", test, ret);
25
26     printf("Returned value: %ld\n", ret);
27
28     return 0;
29 }
30
```

Definition of syscall

Our syscall number

Making a call to the system call.

Notes

- Adding a new system call requires recompiling and reinstalling the whole kernel.
- However, a module can be added at any time without kernel reinstallation.
- That's why, keep the system call definition function (in our case, `test_call.c`) really simple and compile only once. Ideally, this function only creates the `sys_call` pointer and calls it.
- Implement the actual system call handler functions as a module (in our case, `syscallModule.c`). You can compile it as many times as you want.